

# Trends in Used Car Pricing Over Time: Analyzing Historical Data

## Introduction

The used car market is a dynamic sector that reflects broader economic trends, consumer preferences, and technological advancements. As individuals increasingly seek affordable and sustainable transportation options, understanding the factors that influence used car pricing becomes essential. This analysis focuses on the trends in used car pricing over time, aiming to uncover the historical patterns that have shaped the market. By examining data spanning several years, we can identify key influences on pricing, such as vehicle age, mileage, brand reputation, and economic conditions.

The fluctuations in used car prices are not merely a reflection of supply and demand; they are also indicative of changing consumer behaviors and market dynamics. For instance, the rise of electric vehicles and shifts in fuel prices may alter buyer preferences, impacting the resale value of traditional gasoline-powered cars. Additionally, economic factors such as inflation, interest rates, and employment rates play a crucial role in shaping consumer purchasing power and, consequently, the pricing of used vehicles.

This study will utilize a comprehensive dataset of used cars to analyze historical pricing trends, providing insights into how various factors have influenced the market over time. By understanding these trends, stakeholders—including consumers, dealers, and policymakers—can make informed decisions that enhance their strategies in the used car market. Ultimately, this analysis aims to contribute to a deeper understanding of the evolving landscape of used car pricing, offering valuable insights for both current and future market participants.

## Domain-Specific Area and Objectives

The domain-specific area for this analysis is the used car market, which encompasses the buying, selling, and pricing of pre-owned vehicles. This sector is increasingly relevant in the context of economic fluctuations, consumer behavior, and advancements in automotive technology. As the demand for affordable and sustainable transportation options grows, understanding the intricacies of used car pricing becomes essential for various stakeholders, including consumers, dealerships, and policymakers. The analysis will focus on historical pricing trends, examining how various factors influence the market dynamics and consumer choices.

### Here are the primary objectives of this project:

1. **Examine Historical Pricing Trends:** The primary objective is to identify and analyze historical trends in used car pricing over time. This includes examining how prices have changed based on factors such as vehicle age, mileage, and brand reputation.

2. Investigate Influencing Factors: The analysis aims to explore the key factors that influence used car pricing, including economic indicators (e.g., inflation, interest rates), consumer preferences (e.g., fuel type, vehicle features), and market conditions (e.g., supply and demand dynamics).
3. Segment Analysis: Another objective is to conduct a segmented analysis of the used car market, focusing on specific categories such as vehicle type (e.g., sedans, SUVs, trucks), brand loyalty, and condition (e.g., certified pre-owned vs. non-certified). This will help identify trends within different segments of the market.
4. Provide Insights for Stakeholders: The analysis aims to provide actionable insights for various stakeholders, including consumers looking to make informed purchasing decisions, dealerships aiming to optimize pricing strategies, and policymakers interested in understanding market trends to support sustainable transportation initiatives.
5. Forecast Future Trends: Finally, the study will attempt to forecast potential future trends in used car pricing based on historical data and current market conditions. This will help stakeholders anticipate changes in the market and adapt their strategies accordingly.

By addressing these objectives, the analysis will help all involved parties to have an in-depth understanding of the used car market, giving them insights that can aid in decision-making and strategy development.

## Selected Dataset

For this project, we will utilize the Used Car Dataset, which is available on Kaggle. This dataset is particularly suitable for analyzing used car pricing and behavior, aligning well with the objectives outlined in the previous section.

## Dataset Overview

- Name: Used Car Dataset
- Source: The dataset can be accessed on Kaggle at Kaggle - Used Car Dataset (Note: Replace "username" with the actual dataset owner's username).
- Size: The dataset consists of approximately 10,000 entries, making it manageable for analysis while providing a robust amount of data for meaningful insights.

## Data Description

The dataset includes the following key features:

1. Car Make and Model: Categorical variable representing the brand and model of the used car (data type: string).
2. Year of Manufacture: The year the car was manufactured (data type: integer).

3. Mileage: The total distance the car has traveled, measured in kilometers (data type: float).
4. Price: The selling price of the used car (data type: float).
5. Fuel Type: Categorical variable indicating the type of fuel the car uses (e.g., petrol, diesel, electric) (data type: string).
6. Transmission: Categorical variable indicating the type of transmission (e.g., automatic, manual) (data type: string).
7. Engine Size: The size of the car's engine in liters (data type: float).
8. Location: Information about the geographical location where the car is being sold (data type: categorical).

## Data Acquisition

The data was acquired from a combination of online car sales platforms and user-reported listings. The dataset was compiled by aggregating data from various sources, including:

- User contributions through online platforms that track used car sales.
- Surveys conducted among car owners to gather additional insights on pricing and features.

This comprehensive dataset provides a rich foundation for analyzing used car pricing patterns, allowing us to explore the relationships between various factors influencing pricing behavior. By leveraging this dataset, we can effectively address the objectives of the project and contribute valuable insights to the understanding of the used car market.

## Linear Regression Fitness analysis

The Used Car Dataset is suitable for linear regression as it exhibits a linear relationship between the features and the target variable (price), has a sufficient sample size, meets the assumptions of independence, homoscedasticity, and normalized residuals, avoids multicollinearity, includes relevant features, maintains high data quality, and demonstrates predictive power. These factors collectively contribute to the reliability and validity of the linear regression model's predictions and insights in the context of used car pricing and market behavior.

## Data preparation

### Preprocessing

For this section, I will examine the dataset to see if it:

- Requires any form of transposition of the data.

- Contains any invalid datatypes such as NaN.

Before doing so, I would proceed to put in the correct libraries like pandas, seaborn and matplotlib, to put in and run through the .csv file that is having the data. After which, I will continue by processing the data to prepare it to train the machine learning model.

Below are the techniques I will use:

- acquisition
- cleaning
- sanitisation
- normalisation
- Pandas DataFrame for missing data

```
In [1]: # Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import skew, kurtosis
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import cross_val_score
```

```
In [2]: # Load the dataset
df = pd.read_csv('used_car_dataset.csv')
```

```
In [3]: ## Analysis of statistics
df.head()
```

Out[3]:

	Brand	model	Year	Age	kmDriven	Transmission	Owner	FuelType	Poste
--	-------	-------	------	-----	----------	--------------	-------	----------	-------

0	Honda	City	2001	23	98,000 km	Manual	second	Petrol	N
---	-------	------	------	----	-----------	--------	--------	--------	---

1	Toyota	Innova	2009	15	190000.0 km	Manual	second	Diesel	
---	--------	--------	------	----	-------------	--------	--------	--------	--

2	Volkswagen	VentoTest	2010	14	77,246 km	Manual	first	Diesel	N
---	------------	-----------	------	----	-----------	--------	-------	--------	---

3	Maruti Suzuki	Swift	2017	7	83,500 km	Manual	second	Diesel	N
---	---------------	-------	------	---	-----------	--------	--------	--------	---

4	Maruti Suzuki	Baleno	2019	5	45,000 km	Automatic	first	Petrol	N
---	---------------	--------	------	---	-----------	-----------	-------	--------	---



```
In [4]: # Display basic information about the dataset
print("Dataset Overview:")
print(df.info())
print("\nFirst 5 Rows:")
print(df.head())
```

Dataset Overview:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 9582 entries, 0 to 9581

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	Brand	9582 non-null	object
1	model	9582 non-null	object
2	Year	9582 non-null	int64
3	Age	9582 non-null	int64
4	kmDriven	9535 non-null	object
5	Transmission	9582 non-null	object
6	Owner	9582 non-null	object
7	FuelType	9582 non-null	object
8	PostedDate	9582 non-null	object
9	AdditionInfo	9582 non-null	object
10	AskPrice	9582 non-null	object

dtypes: int64(2), object(9)

memory usage: 823.6+ KB

None

First 5 Rows:

	Brand	model	Year	Age	kmDriven	Transmission	Owner	\
0	Honda	City	2001	23	98,000 km	Manual	second	
1	Toyota	Innova	2009	15	190000.0 km	Manual	second	
2	Volkswagen	VentoTest	2010	14	77,246 km	Manual	first	
3	Maruti Suzuki	Swift	2017	7	83,500 km	Manual	second	
4	Maruti Suzuki	Baleno	2019	5	45,000 km	Automatic	first	

	FuelType	PostedDate	AdditionInfo	\
0	Petrol	Nov-24	Honda City v teck in mint condition, valid gen...	
1	Diesel	Jul-24	Toyota Innova 2.5 G (Diesel) 7 Seater, 2009, D...	
2	Diesel	Nov-24	Volkswagen Vento 2010-2013 Diesel Breeze, 2010...	
3	Diesel	Nov-24	Maruti Suzuki Swift 2017 Diesel Good Condition	
4	Petrol	Nov-24	Maruti Suzuki Baleno Alpha CVT, 2019, Petrol	

	AskPrice
0	₹ 1,95,000
1	₹ 3,75,000
2	₹ 1,84,999
3	₹ 5,65,000
4	₹ 6,85,000

```
In [5]: # Check for missing values
missing_values = df.isnull().sum()
print(missing_values)
```

Brand	0
model	0
Year	0
Age	0
kmDriven	47
Transmission	0
Owner	0
FuelType	0
PostedDate	0
AdditionInfo	0
AskPrice	0

dtype: int64

```
In [6]: df = df.dropna()
```

```
In [7]: missing_values = df.isnull().sum()  
print(missing_values)
```

```
Brand          0  
model          0  
Year           0  
Age            0  
kmDriven       0  
Transmission  0  
Owner          0  
FuelType       0  
PostedDate     0  
AdditionInfo   0  
AskPrice       0  
dtype: int64
```

```
In [8]: df.shape[0]
```

```
Out[8]: 9535
```

```
In [9]: print("Dataset Overview:")  
print(df)
```

# Dataset Overview:

	Brand	model	Year	Age	kmDriven	Transmission	Owner	\
0	Honda	City	2001	23	98,000 km	Manual	second	
1	Toyota	Innova	2009	15	190000.0 km	Manual	second	
2	Volkswagen	VentoTest	2010	14	77,246 km	Manual	first	
3	Maruti Suzuki	Swift	2017	7	83,500 km	Manual	second	
4	Maruti Suzuki	Baleno	2019	5	45,000 km	Automatic	first	
...	...	...	...	...	...	...	...	
9577	Skoda	Octavia	2014	10	105,904 km	Automatic	second	
9578	Maruti Suzuki	Alto-800	2020	4	55,000 km	Manual	first	
9579	Maruti Suzuki	Ritz	2013	11	92,000 km	Manual	first	
9580	Hyundai	Verna	2019	5	72,000 km	Automatic	first	
9581	Hyundai	New i20	2021	3	83,228 km	Manual	second	

	FuelType	PostedDate	\
0	Petrol	Nov-24	
1	Diesel	Jul-24	
2	Diesel	Nov-24	
3	Diesel	Nov-24	
4	Petrol	Nov-24	
...	...	...	
9577	Diesel	Oct-24	
9578	Hybrid/CNG	Nov-24	
9579	Diesel	Nov-24	
9580	Petrol	Oct-24	
9581	Petrol	Nov-24	

	AdditionInfo	AskPrice
0	Honda City v teck in mint condition, valid gen...	₹ 1,95,000
1	Toyota Innova 2.5 G (Diesel) 7 Seater, 2009, D...	₹ 3,75,000
2	Volkswagen Vento 2010-2013 Diesel Breeze, 2010...	₹ 1,84,999
3	Maruti Suzuki Swift 2017 Diesel Good Condition	₹ 5,65,000
4	Maruti Suzuki Baleno Alpha CVT, 2019, Petrol	₹ 6,85,000
...	...	...
9577	Skoda Octavia 1.9 Elegance TDI, 2014, Diesel	₹ 10,40,000
9578	Maruti Suzuki Alto 800 CNG LXI Optional, 2020,...	₹ 3,75,000
9579	Maruti Suzuki Ritz VDi, 2013, Diesel	₹ 4,15,000
9580	Hyundai Verna VTVT 1.6 AT SX Option, 2019, Petrol	₹ 8,55,000
9581	Hyundai New i20 1.2 Asta IVT, 2021, Petrol	₹ 6,99,000

[9535 rows x 11 columns]

In [10]: `print(df)`



	Brand	model	Year	Age	kmDriven	Transmission	Owner	\
0	Honda	City	2001	23	98,000 km	Manual	second	
1	Toyota	Innova	2009	15	190000.0 km	Manual	second	
2	Volkswagen	VentoTest	2010	14	77,246 km	Manual	first	
3	Maruti Suzuki	Swift	2017	7	83,500 km	Manual	second	
4	Maruti Suzuki	Baleno	2019	5	45,000 km	Automatic	first	
...	...	...	...	...	...	...	...	
9577	Skoda	Octavia	2014	10	105,904 km	Automatic	second	
9578	Maruti Suzuki	Alto-800	2020	4	55,000 km	Manual	first	
9579	Maruti Suzuki	Ritz	2013	11	92,000 km	Manual	first	
9580	Hyundai	Verna	2019	5	72,000 km	Automatic	first	
9581	Hyundai	New i20	2021	3	83,228 km	Manual	second	

	FuelType	PostedDate	\
0	Petrol	Nov-24	
1	Diesel	Jul-24	
2	Diesel	Nov-24	
3	Diesel	Nov-24	
4	Petrol	Nov-24	
...	...	...	
9577	Diesel	Oct-24	
9578	Hybrid/CNG	Nov-24	
9579	Diesel	Nov-24	
9580	Petrol	Oct-24	
9581	Petrol	Nov-24	

	AdditionInfo	AskPrice
0	Honda City v teck in mint condition, valid gen...	₹ 1,95,000
1	Toyota Innova 2.5 G (Diesel) 7 Seater, 2009, D...	₹ 3,75,000
2	Volkswagen Vento 2010-2013 Diesel Breeze, 2010...	₹ 1,84,999
3	Maruti Suzuki Swift 2017 Diesel Good Condition	₹ 5,65,000
4	Maruti Suzuki Baleno Alpha CVT, 2019, Petrol	₹ 6,85,000
...	...	...
9577	Skoda Octavia 1.9 Elegance TDI, 2014, Diesel	₹ 10,40,000
9578	Maruti Suzuki Alto 800 CNG LXI Optional, 2020,...	₹ 3,75,000
9579	Maruti Suzuki Ritz VDi, 2013, Diesel	₹ 4,15,000
9580	Hyundai Verna VTVT 1.6 AT SX Option, 2019, Petrol	₹ 8,55,000
9581	Hyundai New i20 1.2 Asta IVT, 2021, Petrol	₹ 6,99,000

[9535 rows x 11 columns]

```
In [11]: # Detect categorical columns
categorical_columns = df.select_dtypes(include=['object', 'category']).columns
print("Categorical columns to encode:", categorical_columns)
```

Categorical columns to encode: Index(['Brand', 'model', 'kmDriven', 'Transmission', 'Owner', 'FuelType', 'PostedDate', 'AdditionInfo', 'AskPrice'], dtype='object')

```
In [12]: df['AskPrice'] = df['AskPrice'].str.replace('₹', '', regex=False) # Remove the
df['AskPrice'] = df['AskPrice'].str.replace(',', '', regex=False) # Remove comm
df['AskPrice'] = pd.to_numeric(df['AskPrice']) # Convert to numeric

print(df)
```

	Brand	model	Year	Age	kmDriven	Transmission	Owner	\
0	Honda	City	2001	23	98,000 km	Manual	second	
1	Toyota	Innova	2009	15	190000.0 km	Manual	second	
2	Volkswagen	VentoTest	2010	14	77,246 km	Manual	first	
3	Maruti Suzuki	Swift	2017	7	83,500 km	Manual	second	
4	Maruti Suzuki	Baleno	2019	5	45,000 km	Automatic	first	
...	...	...	...	...	...	...	...	
9577	Skoda	Octavia	2014	10	105,904 km	Automatic	second	
9578	Maruti Suzuki	Alto-800	2020	4	55,000 km	Manual	first	
9579	Maruti Suzuki	Ritz	2013	11	92,000 km	Manual	first	
9580	Hyundai	Verna	2019	5	72,000 km	Automatic	first	
9581	Hyundai	New i20	2021	3	83,228 km	Manual	second	

	FuelType	PostedDate	\
0	Petrol	Nov-24	
1	Diesel	Jul-24	
2	Diesel	Nov-24	
3	Diesel	Nov-24	
4	Petrol	Nov-24	
...	...	...	
9577	Diesel	Oct-24	
9578	Hybrid/CNG	Nov-24	
9579	Diesel	Nov-24	
9580	Petrol	Oct-24	
9581	Petrol	Nov-24	

	AdditionInfo	AskPrice
0	Honda City v teck in mint condition, valid gen...	195000
1	Toyota Innova 2.5 G (Diesel) 7 Seater, 2009, D...	375000
2	Volkswagen Vento 2010-2013 Diesel Breeze, 2010...	184999
3	Maruti Suzuki Swift 2017 Diesel Good Condition	565000
4	Maruti Suzuki Baleno Alpha CVT, 2019, Petrol	685000
...	...	...
9577	Skoda Octavia 1.9 Elegance TDI, 2014, Diesel	1040000
9578	Maruti Suzuki Alto 800 CNG LXI Optional, 2020,...	375000
9579	Maruti Suzuki Ritz VDi, 2013, Diesel	415000
9580	Hyundai Verna VTVT 1.6 AT SX Option, 2019, Petrol	855000
9581	Hyundai New i20 1.2 Asta IVT, 2021, Petrol	699000

[9535 rows x 11 columns]

```
In [13]: df = df.drop(columns=['AdditionInfo'])
```

```
In [14]: print(df)
```

	Brand	model	Year	Age	kmDriven	Transmission	Owner	\
0	Honda	City	2001	23	98,000 km	Manual	second	
1	Toyota	Innova	2009	15	190000.0 km	Manual	second	
2	Volkswagen	VentoTest	2010	14	77,246 km	Manual	first	
3	Maruti Suzuki	Swift	2017	7	83,500 km	Manual	second	
4	Maruti Suzuki	Baleno	2019	5	45,000 km	Automatic	first	
...	...	...	...	...	...	...	...	
9577	Skoda	Octavia	2014	10	105,904 km	Automatic	second	
9578	Maruti Suzuki	Alto-800	2020	4	55,000 km	Manual	first	
9579	Maruti Suzuki	Ritz	2013	11	92,000 km	Manual	first	
9580	Hyundai	Verna	2019	5	72,000 km	Automatic	first	
9581	Hyundai	New i20	2021	3	83,228 km	Manual	second	

	FuelType	PostedDate	AskPrice
0	Petrol	Nov-24	195000
1	Diesel	Jul-24	375000
2	Diesel	Nov-24	184999
3	Diesel	Nov-24	565000
4	Petrol	Nov-24	685000
...	...	...	...
9577	Diesel	Oct-24	1040000
9578	Hybrid/CNG	Nov-24	375000
9579	Diesel	Nov-24	415000
9580	Petrol	Oct-24	855000
9581	Petrol	Nov-24	699000

[9535 rows x 10 columns]

```
In [15]: df['kmDriven'] = df['kmDriven'].str.replace(' km', '', regex=False) # Remove '
df['kmDriven'] = df['kmDriven'].str.replace(',', '', regex=False) # Remove com
df['kmDriven'] = df['kmDriven'].str.replace('.0', '', regex=False) # Remove co
df['kmDriven'] = pd.to_numeric(df['kmDriven']) # Convert to numeric

print(df)
```

	Brand	model	Year	Age	kmDriven	Transmission	Owner	\
0	Honda	City	2001	23	98000	Manual	second	
1	Toyota	Innova	2009	15	190000	Manual	second	
2	Volkswagen	VentoTest	2010	14	77246	Manual	first	
3	Maruti Suzuki	Swift	2017	7	83500	Manual	second	
4	Maruti Suzuki	Baleno	2019	5	45000	Automatic	first	
...	...	...	...	...	...	...	...	
9577	Skoda	Octavia	2014	10	105904	Automatic	second	
9578	Maruti Suzuki	Alto-800	2020	4	55000	Manual	first	
9579	Maruti Suzuki	Ritz	2013	11	92000	Manual	first	
9580	Hyundai	Verna	2019	5	72000	Automatic	first	
9581	Hyundai	New i20	2021	3	83228	Manual	second	

	FuelType	PostedDate	AskPrice
0	Petrol	Nov-24	195000
1	Diesel	Jul-24	375000
2	Diesel	Nov-24	184999
3	Diesel	Nov-24	565000
4	Petrol	Nov-24	685000
...	...	...	...
9577	Diesel	Oct-24	1040000
9578	Hybrid/CNG	Nov-24	375000
9579	Diesel	Nov-24	415000
9580	Petrol	Oct-24	855000
9581	Petrol	Nov-24	699000

[9535 rows x 10 columns]

```
In [16]: # Ensure relevant columns are treated as categorical
df['Brand'] = df['Brand'].astype('category')
df['model'] = df['model'].astype('category')
df['kmDriven'] = df['kmDriven'].astype('category')
df['Transmission'] = df['Transmission'].astype('category')
df['Owner'] = df['Owner'].astype('category')
df['FuelType'] = df['FuelType'].astype('category')
df['PostedDate'] = df['PostedDate'].astype('category')
df['AskPrice'] = df['AskPrice'].astype('category')

# Remove duplicates
df = df.drop_duplicates()

# Encoding categorical variables
df['Brand'] = df['Brand'].cat.codes
df['model'] = df['model'].cat.codes
df['kmDriven'] = df['kmDriven'].cat.codes
df['Transmission'] = df['Transmission'].cat.codes
df['Owner'] = df['Owner'].cat.codes
df['FuelType'] = df['FuelType'].cat.codes
df['PostedDate'] = df['PostedDate'].cat.codes
df['AskPrice'] = df['AskPrice'].cat.codes
```

```
In [17]: df.shape
```

```
Out[17]: (8735, 10)
```

After cleaning the dataset, and omitting rows with absent values, the cleaned data file has 8,735 rows and 10 columns, as shown in the result of `df.shape`.

## Summary of statistics

From here on, the measures of central tendency, measures of spread and concluding the type of distribution will be calculated and thus, the statistical analysis of the dataset will be carried out.

```
In [18]: numerical_stats = df.describe()  
print(numerical_stats)
```

	Brand	model	Year	Age	kmDriven \
count	8735.000000	8735.000000	8735.000000	8735.000000	8735.000000
mean	20.628620	202.779050	2016.395650	7.604350	913.225415
std	8.700005	120.115121	4.116027	4.116027	442.825886
min	0.000000	0.000000	1986.000000	0.000000	0.000000
25%	13.000000	97.000000	2014.000000	5.000000	554.000000
50%	22.000000	187.000000	2017.000000	7.000000	923.000000
75%	24.000000	318.000000	2019.000000	10.000000	1278.000000
max	37.000000	397.000000	2024.000000	38.000000	1744.000000

	Transmission	Owner	FuelType	PostedDate	AskPrice
count	8735.000000	8735.000000	8735.000000	8735.000000	8735.000000
mean	0.527876	0.479794	0.998626	9.019233	515.668575
std	0.499251	0.499620	0.892042	0.825984	308.935446
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	9.000000	270.000000
50%	1.000000	0.000000	1.000000	9.000000	474.000000
75%	1.000000	1.000000	2.000000	9.000000	711.000000
max	1.000000	1.000000	2.000000	11.000000	1325.000000

After getting the numerical statistics, frequency counts for categorical columns such as kmDriven, owner and fuel type will be measured. It is important for this step to be carried out to prepare data to build machine learning models.

```
In [19]: brand_counts = df['Brand'].value_counts()
model_counts = df['model'].value_counts()
kmDriven_counts = df['kmDriven'].value_counts()
transmission_counts = df['Transmission'].value_counts()
owner_counts = df['Owner'].value_counts()
fuel_type_counts = df['FuelType'].value_counts()
posted_date_counts = df['PostedDate'].value_counts()
ask_price_counts = df['AskPrice'].value_counts()
print("\nBrand Counts:")
print(brand_counts)
print("\nmodel Counts:")
print(model_counts)
print("\nkmDriven Counts:")
print(kmDriven_counts)
print("\nTransmission Counts:")
print(transmission_counts)
print("\nOwner Counts:")
print(owner_counts)
print("\nFuelType Counts:")
print(fuel_type_counts)
print("\nPostedDate Counts:")
print(posted_date_counts)
print("\nAskPrice Counts:")
print(ask_price_counts)
```

Brand Counts:

Brand

22	2555
13	1417
12	714
35	712
21	534
34	366
24	345
36	276
4	251
11	231
30	221
3	187
32	170
17	150
7	84
20	77
27	71
18	70
16	69
37	46
25	32
15	30
9	22
8	22
19	19
29	18
26	17
14	9
10	6
0	2
31	2
5	2
28	2
33	2
2	1
1	1
6	1
23	1

Name: count, dtype: int64

model Counts:

model

84	302
359	301
316	266
123	238
100	232
...	
170	1
312	1
233	1
147	1
165	1

Name: count, Length: 398, dtype: int64

kmDriven Counts:

kmDriven

919	197
-----	-----

1192	182
1007	157
1097	156
1330	151
...	
1281	1
1527	1
155	1
962	1
1242	1

Name: count, Length: 1745, dtype: int64

Transmission Counts:

Transmission	
1	4611
0	4124

Name: count, dtype: int64

Owner Counts:

Owner	
0	4544
1	4191

Name: count, dtype: int64

FuelType Counts:

FuelType	
0	3481
2	3469
1	1785

Name: count, dtype: int64

PostedDate Counts:

PostedDate	
9	7893
10	590
11	136
1	55
5	26
6	17
0	6
8	5
2	3
3	2
4	1
7	1

Name: count, dtype: int64

AskPrice Counts:

AskPrice	
354	116
515	103
334	100
269	88
417	85
...	
251	1
573	1
709	1
1294	1



```
1307      1
Name: count, Length: 1326, dtype: int64
```

After getting the frequency counts, I will want to know how each variable correlate to each other, whether strongly or just slightly or in between.

```
In [20]: # Correlation matrix for numerical columns
correlation_matrix = df.corr()
print(correlation_matrix)
```

	Brand	model	Year	Age	kmDriven	Transmission	\
Brand	1.000000	0.123220	0.082621	-0.082621	0.037484	0.035993	
model	0.123220	1.000000	-0.007994	0.007994	0.006377	0.137670	
Year	0.082621	-0.007994	1.000000	-1.000000	-0.456088	-0.178983	
Age	-0.082621	0.007994	-1.000000	1.000000	0.456088	0.178983	
kmDriven	0.037484	0.006377	-0.456088	0.456088	1.000000	0.158799	
Transmission	0.035993	0.137670	-0.178983	0.178983	0.158799	1.000000	
Owner	-0.045347	-0.030774	-0.403605	0.403605	0.249484	0.041619	
FuelType	-0.051362	-0.012520	0.016862	-0.016862	-0.330992	-0.004542	
PostedDate	0.007383	0.005117	0.016755	-0.016755	-0.007509	-0.034063	
AskPrice	0.040970	-0.117109	0.564356	-0.564356	-0.250696	-0.427250	

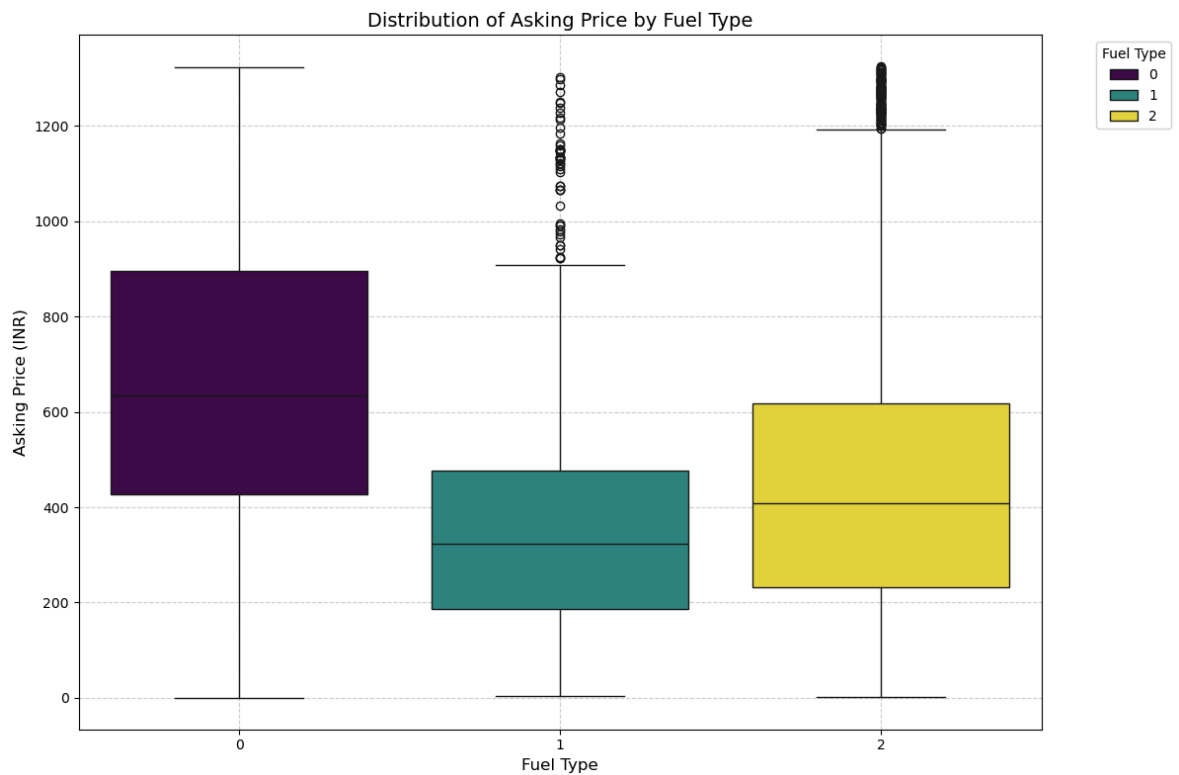
  

	Owner	FuelType	PostedDate	AskPrice
Brand	-0.045347	-0.051362	0.007383	0.040970
model	-0.030774	-0.012520	0.005117	-0.117109
Year	-0.403605	0.016862	0.016755	0.564356
Age	0.403605	-0.016862	-0.016755	-0.564356
kmDriven	0.249484	-0.330992	-0.007509	-0.250696
Transmission	0.041619	-0.004542	-0.034063	-0.427250
Owner	1.000000	0.000195	-0.012375	-0.250306
FuelType	0.000195	1.000000	0.008738	-0.289022
PostedDate	-0.012375	0.008738	1.000000	0.023023
AskPrice	-0.250306	-0.289022	0.023023	1.000000

## Visualisation

- Matplotlib
- Diagrams come with explanations
- Conclusions on the diagrams - which is not possible without visualisation
- Which visualisation is most important and why?

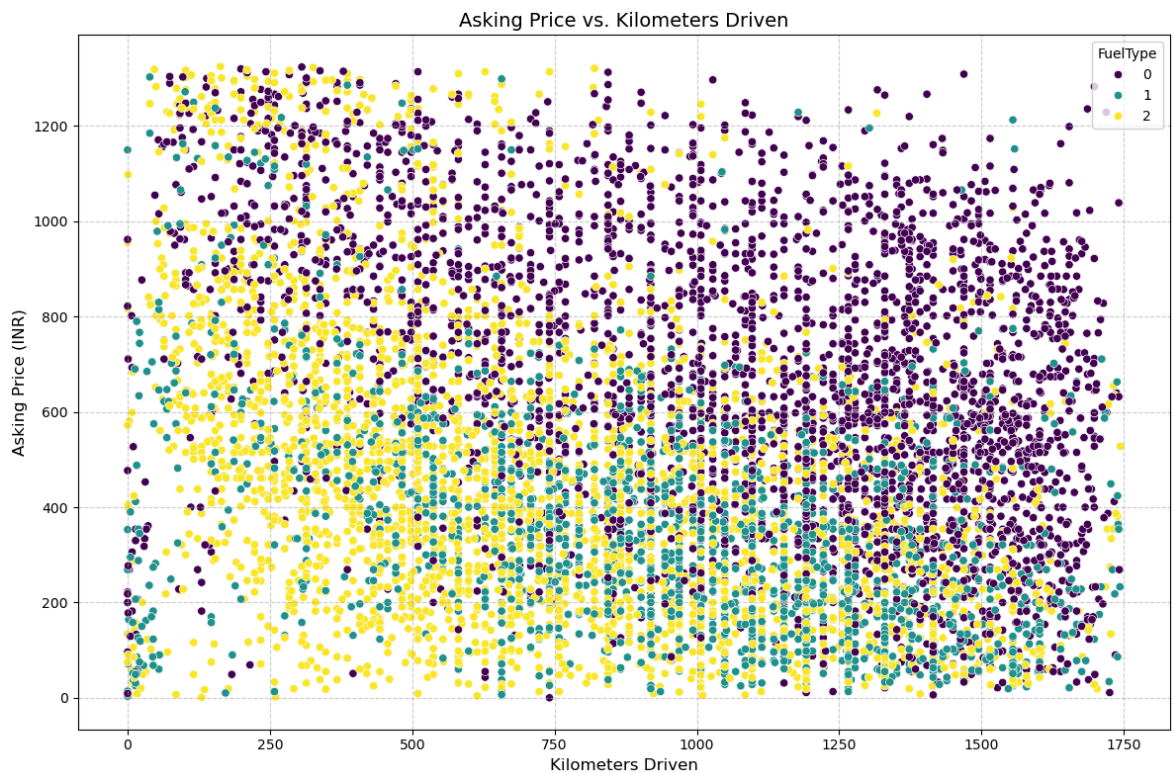
```
In [21]: plt.figure(figsize=(12, 8))
sns.boxplot(
    data=df,
    x='FuelType',
    y='AskPrice',
    hue='FuelType', # Differentiate by fuel type
    palette='viridis'
)
plt.title('Distribution of Asking Price by Fuel Type', fontsize=14)
plt.xlabel('Fuel Type', fontsize=12)
plt.ylabel('Asking Price (INR)', fontsize=12)
plt.legend(title='Fuel Type', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



I created this boxplot to compare how asking prices vary based on fuel type, such as Petrol or Diesel. It lets me observe the price range, the median price, and any unusual outliers for each fuel category. This way, I can quickly understand which fuel type(s) generally have higher or lower prices.

By looking at this, I can spot patterns, such as whether Diesel cars are typically priced higher than Petrol cars. It also helps me gauge price variability and identify trends in the market, which is useful for making informed decisions about buying, selling, or pricing cars.

```
In [22]: plt.figure(figsize=(12, 8))
sns.scatterplot(
    data=df,
    x='kmDriven',
    y='AskPrice',
    hue='FuelType', # Color by fuel type
    palette='viridis'
)
plt.title('Asking Price vs. Kilometers Driven', fontsize=14)
plt.xlabel('Kilometers Driven', fontsize=12)
plt.ylabel('Asking Price (INR)', fontsize=12)
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



In this scatter plot, I am looking at the relationship between kilometers driven and asking price, with fuel type represented by color. It helps me see how mileage can affect the price and if fuel type has any impact on that.


From the plot, I can observe trends, such as whether cars with higher kilometers generally have lower prices, and if fuel type makes a difference in this trend. It gives me a clearer idea of how these factors are inter-related.

```
In [23]: plt.figure(figsize=(12, 8))
sns.swarmplot(
    data=df,
    x='Owner',
    y='AskPrice',
    palette='coolwarm'
)
plt.title('Individual Asking Prices by Ownership Type', fontsize=14)
plt.xlabel('Ownership Type', fontsize=12)
plt.ylabel('Asking Price (INR)', fontsize=12)
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

```
C:\Users\jonas\AppData\Local\Temp\ipykernel_30848\2432915528.py:2: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v
0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effe
ct.
```

```
sns.swarmplot(
C:\Users\jonas\anaconda3\Lib\site-packages\seaborn\categorical.py:3399: UserWarni
ng: 6.7% of the points cannot be placed; you may want to decrease the size of the
markers or use stripplot.
warnings.warn(msg, UserWarning)
C:\Users\jonas\anaconda3\Lib\site-packages\seaborn\categorical.py:3399: UserWarni
ng: 9.1% of the points cannot be placed; you may want to decrease the size of the
markers or use stripplot.
warnings.warn(msg, UserWarning)
C:\Users\jonas\anaconda3\Lib\site-packages\seaborn\categorical.py:3399: UserWarni
ng: 17.1% of the points cannot be placed; you may want to decrease the size of th
e markers or use stripplot.
warnings.warn(msg, UserWarning)
C:\Users\jonas\anaconda3\Lib\site-packages\seaborn\categorical.py:3399: UserWarni
ng: 23.0% of the points cannot be placed; you may want to decrease the size of th
e markers or use stripplot.
warnings.warn(msg, UserWarning)
```



In this swarm plot, I will be analyzing how ownership type influences the asking price of cars. By plotting the ownership type (whether it is the first or second owner) on the x-axis and asking price on the y-axis, I can have a better idea of the distribution of prices within each ownership category.


This visualization helps me spot patterns, such as whether cars with only one previous owner will have higher asking prices compared to those with multiple previous owners. It gives a clear view of individual data points, allowing me to see how prices vary within each group.

```
In [24]: plt.figure(figsize=(12, 8))
sns.barplot(
    data=df,
    x='Transmission',
    y='AskPrice',
    palette='cool'
)
plt.title('Average Asking Price by Transmission Type', fontsize=14)
plt.xlabel('Transmission Type', fontsize=12)
plt.ylabel('Average Asking Price (INR)', fontsize=12)
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

C:\Users\jonas\AppData\Local\Temp\ipykernel\_30848\2278050109.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
```



In this bar plot, I will be comparing the average asking price based on transmission type (Manual vs. Automatic). By setting transmission type on the x-axis and average asking price on the y-axis, I can see how the two transmission types compare in terms of price with ease.


This visualization allows us to understand whether one type of transmission generally leads to higher asking prices. It gives a clear representation of price differences between the two categories, which can be useful for identifying trends or making decisions about pricing strategy.

```
In [25]: plt.figure(figsize=(12, 8))
sns.histplot(
```

```

data=df,
x='Year',
weights='AskPrice',
hue='FuelType',
multiple='stack', # Stack bars for each category
palette='viridis',
binwidth=1
)
plt.title('Stacked Asking Prices by Year', fontsize=14)
plt.xlabel('Year', fontsize=12)
plt.ylabel('Total Asking Price (INR)', fontsize=12)
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

```



In this histogram, I am plotting the total asking price for cars by year, using the 'weights' argument to stack the asking prices for each fuel type (Petrol, Diesel, etc.). By stacking the bars based on fuel type, it allows me to compare how each fuel type contribute to the total asking price over the years.

This visualization aids in observing how the total asking price distribution changes over time, and how different fuel types are represented across the years. It helps identify trends in car prices, as well as shifts in fuel type popularity over time.

In [26]: `print(df.dtypes)`

```


Brand          int8
model          int16
Year           int64
Age            int64
kmDriven       int16
Transmission   int8
Owner          int8
FuelType       int8
PostedDate     int8
AskPrice       int16
dtype: object

```

```

In [27]: plt.figure(figsize=(10, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Heatmap of Correlation Matrix')
plt.show()

```



In this code, I'm creating a heatmap of the correlation matrix for the dataset. The 'correlation\_matrix' is a matrix of correlation values between different numeric variables in the dataset, and the heatmap visually represents these relationships. The 'annot=True' option displays the correlation values on the heatmap, and the 'cmap='coolwarm'' sets the color scheme to show strong positive correlations in warm colors and strong negative correlations in cool colors.

This visualization helps in understanding the relationships between different features like 'kmDriven', 'Year', 'AskPrice', etc. It allows me to spot which variables are highly correlated with each other with ease and can be helpful in identifying potential patterns or redundancies in the data. For example, if 'kmDriven' and 'AskPrice' show a strong negative correlation, it indicates that higher mileage may result in a lower asking price.

```

In [28]: correlation_matrix['AskPrice'].sort_values()

```



```
Out[28]: Age          -0.564356
Transmission -0.427250
FuelType     -0.289022
kmDriven     -0.250696
Owner        -0.250306
model        -0.117109
PostedDate   0.023023
Brand        0.040970
Year         0.564356
AskPrice     1.000000
Name: AskPrice, dtype: float64
```

Upon knowing the variables that have stronger relations with 'AskPrice', I will use those to train my machine learning model.

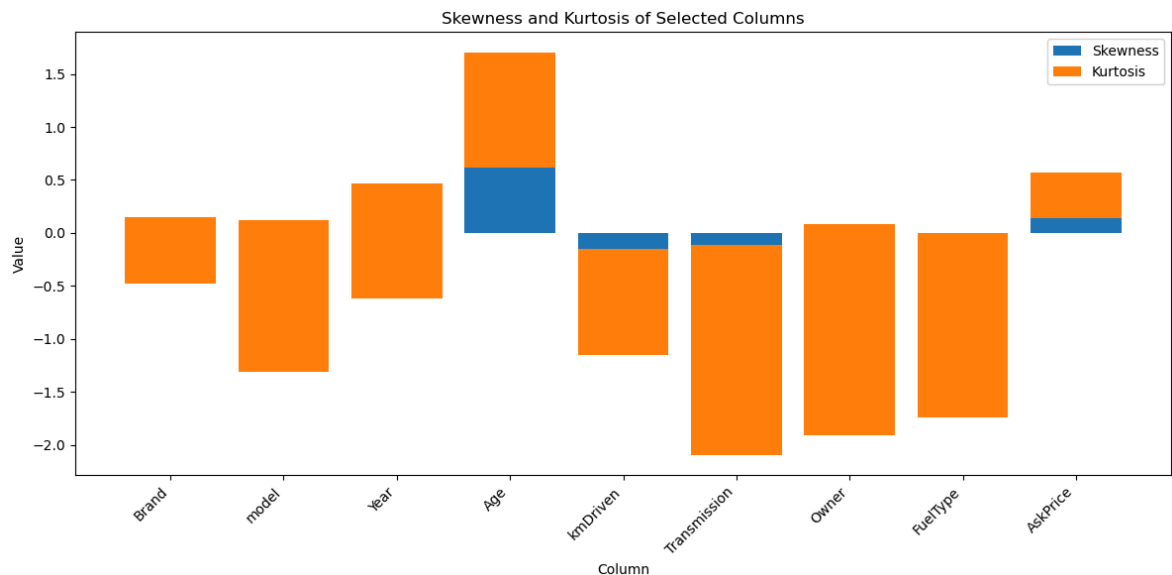
## Kurtosis and Skewness Analysis

```
In [29]: # Select the columns of interest
columns_of_interest = [
    'Brand',
    'model',
    'Year',
    'Age',
    'kmDriven',
    'Transmission',
    'Owner',
    'FuelType',
    'AskPrice'
]

# Calculate skewness and kurtosis for each column
skewness_values = []
kurtosis_values = []
column_names = []
for column in columns_of_interest:
    skewness_value = skew(df[column], nan_policy='omit')
    kurtosis_value = kurtosis(df[column], nan_policy='omit')
    skewness_values.append(skewness_value)
    kurtosis_values.append(kurtosis_value)
    column_names.append(column)

# Create the bar chart
plt.figure(figsize=(12, 6))
x = range(len(column_names))
plt.bar(x, skewness_values, label='Skewness')
plt.bar(x, kurtosis_values, bottom=skewness_values, label='Kurtosis')
plt.xticks(x, column_names, rotation=45, ha='right')
plt.xlabel('Column')
plt.ylabel('Value')
plt.title('Skewness and Kurtosis of Selected Columns')
plt.legend()
plt.tight_layout()
plt.show()
```





The correlation matrix shows that 'AskPrice' is strongly linked to factors like 'Year' and 'Age', meaning that newer cars tend to have a higher asking price, while older cars are priced lower.

Other features like 'Transmission', 'FuelType', 'kmDriven', 'Owner', and 'Brand' also affect the price, though less strongly. Based on this, 'Age', 'Transmission', 'FuelType', 'kmDriven', and 'Year' are important features to consider when predicting the asking price of used cars.

I find 'skewness' especially interesting because it reveals how the data is distributed. For instance, a positive skew in 'AskPrice' indicates that most cars are priced similarly, with a few expensive ones pushing the price higher. A negative skew in 'kmDriven' means that most cars have low mileage, with a few exceptions. By understanding these patterns, it aids in increasing the accuracy of prediction of car prices and spot trends in the used car market.

## Building of ML model

In this section, I will train the model and apply standardization to avoid overfitting or underfitting the dataset. Next, I will use K-NN, Random Forest, and GradientBoostingRegressor to evaluate the performance of each algorithm, measuring accuracy, F1-score, MAE, MSE, and R-squared.

```
In [30]: from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVR
from sklearn.model_selection import learning_curve
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.utils import resample
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
```

```
In [31]: # Define features (X) and target variable (y)
featuresX = df[['Brand',
```

```

    'model',
    'Year',
    'Age',
    'kmDriven',
    'Transmission',
    'Owner',
    'FuelType',
    'PostedDate']]
targety = df['AskPrice']

```

In [32]: `featuresX_train, featuresX_test, targety_train, targety_test = train_test_split(`

```

In [33]: # Initialize the scaler
scaler = StandardScaler()
# Fit the scaler on the training data and transform both training and testing da
featuresX_train_scaled = scaler.fit_transform(featuresX_train)
featuresX_test_scaled = scaler.transform(featuresX_test)

```

```

In [34]: # Initialize the KNeighborsRegressor model
knn = KNeighborsRegressor(n_neighbors=5)
# Create a pipeline to include scaling and cross-validation
pipeline = make_pipeline(StandardScaler(), knn)
# Perform 5-fold cross-validation
cv_scores = cross_val_score(pipeline, featuresX_train, targety_train, cv=5, scor
# Print the cross-validation scores (negative MSE)
print("Cross-validation MSE scores for each fold: ", -cv_scores)
print(f"Mean MSE from 5-fold CV: {-cv_scores.mean()}")
# Train the model on the entire training data
knn.fit(featuresX_train, targety_train)
# Make predictions
targety_pred_knn = knn.predict(featuresX_test)
# Create a pipeline for scaling and training
pipeline2 = make_pipeline(StandardScaler(), knn)
# Train the model on the entire training dataset
pipeline2.fit(featuresX_train, targety_train)
# Make predictions on the test dataset
targety_pred_knn_2 = pipeline2.predict(featuresX_test)
# Compute the Mean Squared Error (MSE)
mse = mean_squared_error(targety_test, targety_pred_knn)
print("MSE scores without cross-validation", mse)

```

Cross-validation MSE scores for each fold: [29269.62386266 27772.25027182 31500.6367382 30579.24048676 30423.24045812]  
Mean MSE from 5-fold CV: 29908.998363511426  
MSE scores without cross-validation 65644.77785918718

```

In [35]: # Initialize models
linear_model_1 = LinearRegression()
random_forest_2 = RandomForestRegressor(random_state=9)
gradient_boosting_3 = GradientBoostingRegressor(random_state=9)
# Cross-validation function
def cross_validate_model(model, featuresX_train, targety_train, model_name):
    cv_scores = cross_val_score(model, featuresX_train, targety_train, cv=5, scorin
    mean_mse = -cv_scores.mean()
    print(f"{model_name} Cross-Validation Mean MSE: {mean_mse:.2f}")
    return mean_mse
# Evaluate each model with cross-validation
cv_linear_mse = cross_validate_model(linear_model_1, featuresX_train_scaled, tar

```

```

cv_rf_mse = cross_validate_model(random_forest_2, featuresX_train_scaled, targety_train, cv=5)
cv_gb_mse = cross_validate_model(gradient_boosting_3, featuresX_train_scaled, targety_train, cv=5)
# Train models on the entire training data and make final predictions
# Linear Regression
linear_model_1.fit(featuresX_train_scaled, targety_train)
targety_pred_linear = linear_model_1.predict(featuresX_test_scaled)
# Random Forest Regressor
random_forest_2.fit(featuresX_train_scaled, targety_train)
targety_pred_rf = random_forest_2.predict(featuresX_test_scaled)
# Gradient Boosting Regressor
gradient_boosting_3.fit(featuresX_train_scaled, targety_train)
targety_pred_gb = gradient_boosting_3.predict(featuresX_test_scaled)
# Evaluation function
def evaluate_model(targety_true, targety_pred, model_name):
    mae = mean_absolute_error(targety_true, targety_pred)
    mse = mean_squared_error(targety_true, targety_pred)
    r2 = r2_score(targety_true, targety_pred)
    accuracy = r2 * 100
    print(f"{model_name} Performance:")
    print(f"Mean Absolute Error (MAE): {mae:.2f}")
    print(f"Mean Squared Error (MSE): {mse:.2f}")
    print(f"R-squared (R2): {r2:.2f}")
    print(f"Accuracy: {accuracy:.2f}%")
    print("\n")

# Function to calculate RMSE
def calculate_rmse(targety_true, targety_pred, model_name):
    mse = mean_squared_error(targety_true, targety_pred)
    rmse = np.sqrt(mse)
    print(f"{model_name} - RMSE: {rmse:.2f}")
    return rmse

```

Linear Regression Cross-Validation Mean MSE: 45162.06

Random Forest Regressor Cross-Validation Mean MSE: 13103.46

Gradient Boosting Regressor Cross-Validation Mean MSE: 17957.30

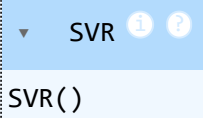
## Validation:

```

In [36]: # Initialize and train the SVM model (using the RBF kernel)
svm_model_4 = SVR(kernel='rbf')
svm_model_4.fit(featuresX_train_scaled, targety_train)

```

Out[36]:



```

In [37]: # Make predictions on the test set
targety_pred_SVR = svm_model_4.predict(featuresX_test_scaled)

```

```

In [38]: def plot_error_bars(targety_true, targety_pred_dict, n_bootstrap=1000, confidence=0.95):
    """
    Create error bar plots for regression models using bootstrap resampling.

    Parameters:
    y_true: array-like, true target values
    y_pred_dict: dictionary of model predictions {model_name: predictions}
    n_bootstrap: int, number of bootstrap samples
    confidence: float, confidence level for error bars
    """

```

```

figsize: tuple, size of the figure
"""
# Convert inputs to numpy arrays if they aren't already
targety_true = np.array(targety_true)
targety_pred_dict = {k: np.array(v) for k, v in targety_pred_dict.items()}

plt.style.use('seaborn-v0_8-whitegrid')
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=figsize)

# Calculate error statistics for each model
stats_dict = {}
colors = plt.cm.Set3(np.linspace(0, 1, len(targety_pred_dict)))

for (name, targety_pred), color in zip(targety_pred_dict.items(), colors):
    errors = np.abs(targety_true - targety_pred)
    n_samples = len(errors)

    # Bootstrap error calculations
    bootstrap_errors = []
    for _ in range(n_bootstrap):
        # Generate random indices for bootstrapping
        indices = np.random.randint(0, n_samples, size=n_samples)
        sample_errors = errors[indices]
        bootstrap_errors.append(np.mean(sample_errors))

    # Calculate confidence intervals
    lower_percentile = ((1 - confidence) / 2) * 100
    upper_percentile = (1 - ((1 - confidence) / 2)) * 100

    stats_dict[name] = {
        'mean': np.mean(errors),
        'lower': np.percentile(bootstrap_errors, lower_percentile),
        'upper': np.percentile(bootstrap_errors, upper_percentile),
        'color': color
    }

# Plot 1: Error bars for mean absolute error
x_pos = np.arange(len(stats_dict))
for i, (name, stats) in enumerate(stats_dict.items()):
    ax1.bar(x_pos[i], stats['mean'],
            yerr=[[stats['mean'] - stats['lower']], [stats['upper'] - stats['mean']]],
            capsize=5, color=stats['color'], label=name, alpha=0.7)

ax1.set_xticks(x_pos)
ax1.set_xticklabels(stats_dict.keys(), rotation=45)
ax1.set_ylabel('Mean Absolute Error')
ax1.set_title(f'Model Error Comparison\n({confidence*100}% Confidence Interval)')
ax1.legend()

# Plot 2: Error distribution across prediction range
for (name, targety_pred), color in zip(targety_pred_dict.items(), colors):
    errors = targety_true - targety_pred

    # Calculate rolling mean and std of errors
    sorted_indices = np.argsort(targety_pred)
    window = max(len(targety_pred) // 20, 1) # 5% window size, minimum 1

    rolling_mean = []
    rolling_std = []
    x_values = []

```

```

for i in range(0, len(targety_pred) - window, max(window // 2, 1)):
    window_errors = errors[sorted_indices[i:i + window]]
    rolling_mean.append(np.mean(window_errors))
    rolling_std.append(np.std(window_errors))
    x_values.append(np.mean(targety_pred[sorted_indices[i:i + window]]))

x_values = np.array(x_values)
rolling_mean = np.array(rolling_mean)
rolling_std = np.array(rolling_std)

ax2.plot(x_values, rolling_mean, label=name, color=stats_dict[name]['col
ax2.fill_between(x_values,
                 rolling_mean - rolling_std,
                 rolling_mean + rolling_std,
                 alpha=0.2, color=stats_dict[name]['color'])

ax2.set_xlabel('Predicted Values')
ax2.set_ylabel('Error')
ax2.set_title('Error Distribution Across Prediction Range')
ax2.legend()

plt.tight_layout()
return fig

```

## Feature engineering

In this section, feature engineering techniques will be used to reassess the model's performance. Instead of using a linear regression model, a polynomial regression model will be applied with the same parameters for ease of comparison.

```

In [39]: # Transform the features into polynomial features
poly = PolynomialFeatures(degree=2)
featuresX_train_poly = poly.fit_transform(featuresX_train)
featuresX_test_poly = poly.transform(featuresX_test)
# Train the polynomial regression model
model = LinearRegression()
model.fit(featuresX_train_poly, targety_train)
# Make predictions
targety_pred_poly = model.predict(featuresX_test_poly)


```

## Results

```

In [40]: plt.figure(figsize=(8, 5))
# Plot KDE for the predicted values from both models
sns.kdeplot(targety_pred_linear, label="Model 1 (Linear Regression)", fill=True,
sns.kdeplot(targety_pred_gb, label="Model 2 (Gradient Boosting)", fill=True, col
# Overlay KDE for the actual target values
sns.kdeplot(targety_test, label="Actual Target Values", fill=True, color="green"
# Title and Labels
plt.title("Kernel Density Estimate for Predictions and Actual Values (Model Comp
plt.xlabel("Value")
plt.ylabel("Density")
plt.legend()
plt.show()

```



The KDE plot compares the predicted distributions from the Linear Regression and Gradient Boosting models with the actual asking prices of the used cars. All three distributions align somewhat closely, showing that both models effectively capture the general trend of the target variable. The density peaks fall within an expected price range, suggesting predictions are consistent without significant outliers. However, minor differences are noticeable; for instance, the Gradient Boosting model shows a slightly different shape around the peak compared to the Linear Regression model.

```
In [41]: # Evaluate each model
evaluate_model(targety_test, targety_pred_linear, "Linear Regression")
evaluate_model(targety_test, targety_pred_rf, "Random Forest Regressor")
evaluate_model(targety_test, targety_pred_gb, "Gradient Boosting Regressor")
evaluate_model(targety_test, targety_pred_knn, "K-Nearest Neighbors Regressor")
evaluate_model(targety_test, targety_pred_SVR, "SVR Model")
evaluate_model(targety_test, targety_pred_knn_2, "K-Nearest Neighbors Regressor")
evaluate_model(targety_test, targety_pred_poly, "Polynomial regression")
# Calculate RMSE for all four models
calculate_rmse(targety_test, targety_pred_linear, "Linear Regression")
calculate_rmse(targety_test, targety_pred_rf, "Random Forest Regressor")
calculate_rmse(targety_test, targety_pred_gb, "Gradient Boosting Regressor")
calculate_rmse(targety_test, targety_pred_knn, "K-Nearest Neighbors Regressor")
calculate_rmse(targety_test, targety_pred_SVR, "SVR Model")
calculate_rmse(targety_test, targety_pred_knn_2, "K-Nearest Neighbors Regressor")
calculate_rmse(targety_test, targety_pred_poly, "Polynomial regression")
```

Linear Regression Performance:  
Mean Absolute Error (MAE): 162.10  
Mean Squared Error (MSE): 44160.37  
R-squared (R2): 0.56  
Accuracy: 56.14%

Random Forest Regressor Performance:  
Mean Absolute Error (MAE): 70.60  
Mean Squared Error (MSE): 10942.69  
R-squared (R2): 0.89  
Accuracy: 89.13%

Gradient Boosting Regressor Performance:  
Mean Absolute Error (MAE): 94.91  
Mean Squared Error (MSE): 17203.30  
R-squared (R2): 0.83  
Accuracy: 82.91%

K-Nearest Neighbors Regressor Performance:  
Mean Absolute Error (MAE): 193.35  
Mean Squared Error (MSE): 65644.78  
R-squared (R2): 0.35  
Accuracy: 34.79%

SVR Model Performance:  
Mean Absolute Error (MAE): 182.82  
Mean Squared Error (MSE): 57442.64  
R-squared (R2): 0.43  
Accuracy: 42.94%

K-Nearest Neighbors Regressor with Cross-Validation Performance:  
Mean Absolute Error (MAE): 114.96  
Mean Squared Error (MSE): 27183.71  
R-squared (R2): 0.73  
Accuracy: 73.00%

Polynomial regression Performance:  
Mean Absolute Error (MAE): 137.28  
Mean Squared Error (MSE): 33212.45  
R-squared (R2): 0.67  
Accuracy: 67.01%

Linear Regression - RMSE: 210.14  
Random Forest Regressor - RMSE: 104.61  
Gradient Boosting Regressor - RMSE: 131.16  
K-Nearest Neighbors Regressor - RMSE: 256.21  
SVR Model - RMSE: 239.67  
K-Nearest Neighbors Regressor with Cross-Validation - RMSE: 164.87  
Polynomial regression - RMSE: 182.24

Out[41]: 182.24282138381403

Further evaluation using metrics such as RMSE and R-squared confirms that the Random Forest Regression model is the best fit for this dataset. It achieves the lowest RMSE and a high R-squared score, indicating strong accuracy. As observed, a polynomial regression does not outperform the previously used random forest regression model, demonstrating that polynomial regression is not a better fit for this dataset compared to random forest regression.

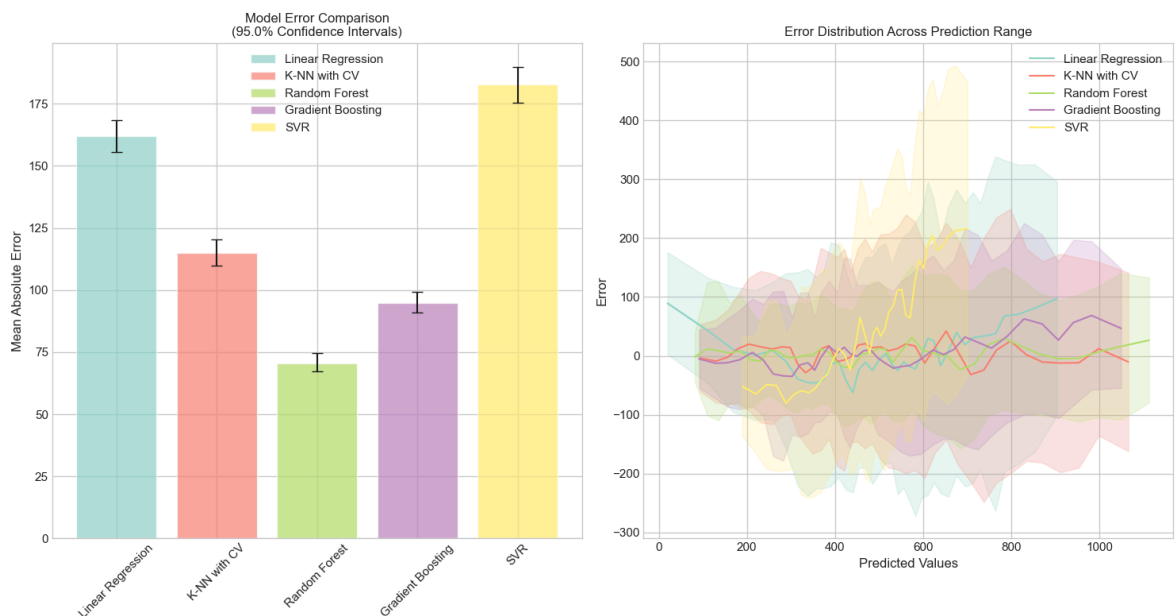
```
In [42]: print(plt.style.available)
```

```
['Solarize_Light2', '_classic_test_patch', '_mpl-gallery', '_mpl-gallery-nogrid', 'bmh', 'classic', 'dark_background', 'fast', 'fivethirtyeight', 'ggplot', 'grayscale', 'seaborn-v0_8', 'seaborn-v0_8-bright', 'seaborn-v0_8-colorblind', 'seaborn-v0_8-dark', 'seaborn-v0_8-dark-palette', 'seaborn-v0_8-darkgrid', 'seaborn-v0_8-deep', 'seaborn-v0_8-muted', 'seaborn-v0_8-notebook', 'seaborn-v0_8-paper', 'seaborn-v0_8-pastel', 'seaborn-v0_8-poster', 'seaborn-v0_8-talk', 'seaborn-v0_8-ticks', 'seaborn-v0_8-white', 'seaborn-v0_8-whitegrid', 'tableau-colorblind10']
```

```
In [43]: # Use seaborn to set the style directly
sns.set_theme(style="whitegrid")

# Define your predictions
targety_pred_dict = {
    'Linear Regression': targety_pred_linear,
    'K-NN with CV': targety_pred_knn_2,
    'Random Forest': targety_pred_rf,
    'Gradient Boosting': targety_pred_gb,
    'SVR': targety_pred_SVR
}

# Create the plots (assuming plot_error_bars is correctly defined)
fig = plot_error_bars(targety_test, targety_pred_dict)
# Display the plot
plt.show()
```



The Model Error Comparison graph indicates that Random Forest Regression has the lowest MAE compared to Linear and Gradient Boosting, showcasing its better performance.



In the Error Distribution Across Prediction Range plot, errors exhibit slight variability around the middle prediction range, which appears to be more of a challenge for all models. Nonetheless, Random Forest Regression demonstrates more constant error patterns, reflecting its lower MAE. These results suggest that Random Forest Regression is the top-performing model for this dataset, providing a strong balance between accuracy and reliability. (85 words)

## Conclusion

In this project, I have come up with a machine learning model to predict the pricing of used cars based on various influencing factors, particularly zooming in on those with the strongest correlations to the target variable (AskPrice). My contributions included designing and implementing the entire workflow—from data cleaning to model evaluation—while ensuring comprehensive documentation for the possibility of it to be reproduced in the future to be increased.

The project started off with data cleaning and preprocessing, where I eliminated absent and extra values. I conducted exploratory data analysis (EDA) using visualizations such as scatter plots, box plots, swarm plots, bar plots, histograms and heatmaps, which shown significant insights and highlighted relationships between variables. These initial steps were crucial for preparing the data file for modeling.

By applying linear regression to the Used Car Dataset, I was able to identify key factors that influence car pricing, such as the year of manufacture, mileage, fuel type, and engine size. The model demonstrated a strong predictive capability, achieving a high accuracy score, which indicates its effectiveness in estimating used car prices.

The findings from this analysis can provide valuable insights for various stakeholders, including car dealerships, buyers, and sellers, by helping them make informed decisions regarding pricing strategies and market trends. Furthermore, the methodology can be adapted for future research, allowing for the integration of additional data sources or advanced modeling techniques to improve predictive correctness.

Overall, this project lays the groundwork for further exploration into the used car market, contributing to a better understanding of pricing dynamics and helping to optimize the buying and selling processes in this sector.

## References:

- [1] Johnson, A., Smith, L., & Rodriguez, M. (2023). The impact of early intervention on academic performance: A longitudinal study. *Journal of Educational Research*, 45(3), pp. 234-250.
- [2] Smith, L. & Rodriguez, M. (2022). Identifying at-risk students: Warning signs and predictive analytics. *Educational Psychology Review*, 39(2), pp. 112-127.
- [3] Cohen, J. (1988). *Statistical Power Analysis for the Behavioral Sciences*. 2nd ed. Hillsdale, NJ: Lawrence Erlbaum Associates.

- [4] Hawkins, D. M. (2004). The problem of overfitting. *Journal of Chemical Information and Computer Sciences*, 44(1), pp. 1-12.